

Regmap - Linux Kernel Subsystems for peripheral devices

Jan Richter-Brockmann, Moritz Schrey

Abstract—The regmap structure is a very beneficial and easy to use implementation for kernel modules. This paper should give an overview of the advantages and disadvantages of regmap. Furthermore, it will follow a short explanation about the implementation of the initialisation and the read and write accesses. The next section will deal with the internal debugging mode of regmap. Finally, it follows a short extract of the *at86rf230* module, where the regmap structure is used.

Keywords—*regmap, SPI, ASoC, debugfs, at86rf230.*

I. INTRODUCTION

MANY Serial Peripheral Interfaces (SPI) and Integrated Circuits (I2C) based devices implement register maps. These register maps are placed on top of the raw wire interface and the implementation is often done in a very similar way. The result is the same code in a lot of drivers of SPI masters.

The idea of the implementation of regmap came from the ALSA System on Chip (ASoC) Layer, which provides a better Advanced Linux Sound Architecture (ALSA) support for embedded system on chip processors and portable audio codecs. However, codec drivers usually connected to the underlaying SoC CPU. Before the ASoC was implemented this led to similar code in different drivers for different platforms [1].

Furthermore, the regmap implementation has more features besides avoiding code doubling. One of these features is to reduce the traffic on the SPI bus. Regmap provides a register cache which contains an image of most registers from the device. Another improvement is the support of a debugging mode.

The next sections will give a detailed explanations about the implementation and functionality of these regmap features.

II. IMPLEMENTATION

A. Configuration and Initialisation

If you want to use the regmap implementation in a driver with a SPI communication you have to call the function *devm_regmap_init_spi(dev, config)*.

This function is used to initialise the regmap structure. The first argument is a pointer to the SPI settings which has to be configured for the connected device. The second parameter is a pointer to a *regmap_config* struct. An extract of this struct is shown in Listing 1 [2]. It contains the settings for the regmap structure which are necessary for its functionality.

Mandatory settings are the number of bits in a register address (*reg_bits*) and the number of bits in a register value (*val_bits*).

Moreover, the read and write flags can be set to define the SPI commands (read, write, ...) as forced by the connected SPI slave. to tell the devices the direction of the SPI communication.

Another important part is the setting of the functions, which return the behaviour of a register, due to four different existing types.

- readable registers
- writeable registers
- volatile registers
- precious registers

The first ones are the read only registers. The second types are writeable registers which are always readable registers too. The next ones are registers, which can change their states during the runtime of the device (volatile registers). The last ones are precious registers, which should not be read outside of a call from the driver [2].

The last configuration that should be discussed is the setting of the cache type. Actually, there exists three possible cache types in the regmap implementation. The first type uses a cache organisation with a red-black tree. The advantage is a balanced binary search tree with a short read access [3]. The second type applies the Lempel-Ziv-Oberhumer (LZO) lossless data compression algorithm [4]. The last type uses a flat cache organisation.

Listing 1. Extract of the struct *regmap_config* definition

```
struct regmap_config {
    const char *name;

    int reg_bits;
    int reg_stride;
    int pad_bits;
    int val_bits;

    bool (*writeable_reg)(struct device *dev, unsigned int reg);
    bool (*readable_reg)(struct device *dev, unsigned int reg);
    bool (*volatile_reg)(struct device *dev, unsigned int reg);
    bool (*precious_reg)(struct device *dev, unsigned int reg);

    // ...

    enum regcache_type cache_type;

    // ...

    u8 read_flag_mask;
    u8 write_flag_mask;

    // ...
};
```

Calling *devm_regmap_init_spi()* causes the execution of the function

regcache_init().

This function is responsible for filling the register cache. This procedure is shown in Figure 1. An if-statement checks if a default register setting is stored in the regmap configuration under *config->reg_defaults*. If this condition is true, the

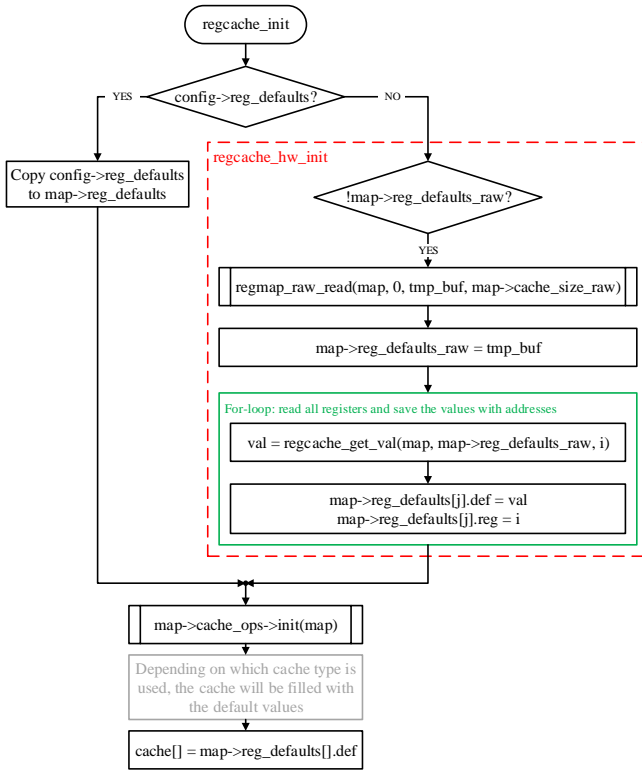


Fig. 1. Initialisation of the register cache

default register settings are copied into the regmap struct¹ (*map->reg_defaults*). If not, the register values have to be read from the device and then stored into the regmap struct. Therefore, the function

regcache_hw_init()

is called. First, it is determined which cache type was selected. With this information the functions for the cache operations can be stored in *map->cache_ops*. Especially this states the read, write and initialisation functions for the cache. If *map->reg_defaults_raw* is empty, the function

regmap_raw_read()

is called. This function reads the raw register values from the device and stores them into the temporary buffer *tmp_buf*. After finishing this reading process the temporary buffer is transferred into *map->reg_defaults_raw*. At this moment, only the register values are stored without the register addresses. A for-loop is going to change this and gets the raw register values with the function *regcache_get_val()* from *map->reg_defaults_raw*. Now the register values can be stored with the register addresses in the struct *map->reg_defaults[]*.

At this point the register values are stored and the cache can be filled. Depending on which cache type was chosen a different initialisation function is called over *map->cache_ops->init()*.

¹The regmap struct contains all settings and functions for regmap. The implementation can be found in [5].

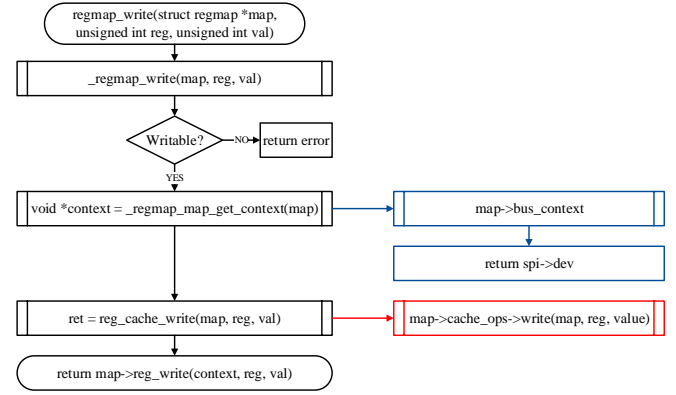


Fig. 2. Write access

B. Write Access

Figure 2 shows the process when a write access is executed. The function

regmap_write()

must be called in the driver. It causes the call of

_regmap_write()

where the real work is done to execute the write access. The first action is to check the register for write permissions. If it is possible to write into the register, the next step is to get the bus context with help of the function

_regmap_map_get_context().

In the case of using SPI communication this function will return the information about the SPI device. After that, the function

reg_cache_write()

is called. Its job is to change the register value in the register cache. Therefore, the function *map->cache_ops->write(map, reg, value)* will be called, which depends on the the cache type that is used. Finally, in the return statement *_regmap_write* calls the function *map->reg_write* to change the desired value on the hardware.

C. Read Access

To get a read access to the regmap structure, the function

regmap_read()

has to be called. The process of a read access is shown in Figure 3. *regmap_read* calls the function

_regmap_read()

which gets the bus context like it is described in the write access. After that, an if-statement checks if a bypass function is enabled. An enabled bypass function means that only the hardware is modified and not the cache, when a write access is executed [5]. However, if the bypass function is disabled, the function

_regcache_read()

is called. When the register is not volatile, the register value can be read from the cache with the function *map->cache_ops->read(map, reg, val)*, which also depends on the chosen cache type. The return value is true and the read access is finished. Nevertheless, it is possible that

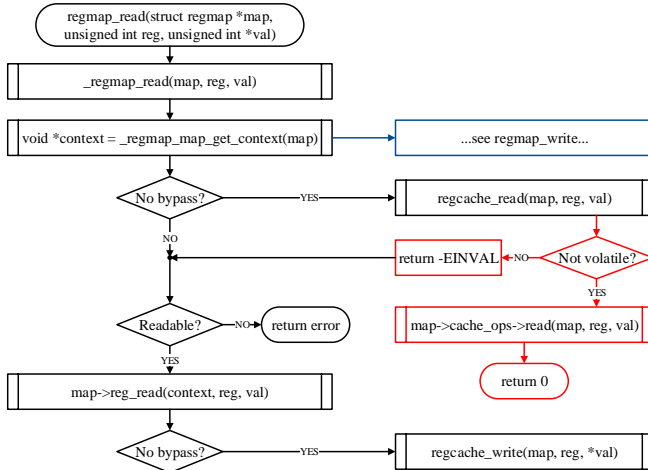


Fig. 3. Read access

the register is a volatile one. With the help of a negative return value the program jumps back into `_regmap_read`. After checking the register for an allowed read access the function `map->reg_read` is called. This function reads the register value from the hardware device. If the bypass function is disabled, the register value is transferred to the regmap cache.

III. DEBUGGING WITH REGMAP

A. debugfs

During kernel programming it is often useful to export debugging information to the user space. One way could be to use `printf()` commands. It often works but if the developer wants to change values from user space, `printf()` is not the solution. To handle this problem, it is common to use virtual file systems like `sysfs`². However, to keep things simple, Greg Kroah-Hartman wrote the special file system `debugfs` [7]. It is a simple-to-use (RAM-based) file system which can read data from the kernel and make it visible to the user space. Furthermore, it is possible to write data from the user space to the kernel. In contrast to the `sysfs` file system, which has the one-value-per-file rule, `debugfs` has no special rules [8].

Usage of `debugfs` requires to include the header file `<linux/debugfs.h>`. Then, at least one directory has to be created to hold a set of `debugfs` files. To create this directory, the function

```
struct dentry *debugfs_create_dir()
```

has to be called. The transferred parameters are `const char *name` and `struct dentry *parent` which leads to the creation of a directory called `name` in the indicated `parent` directory. If the pointer `parent` is NULL, the directory is created in the `debugfs` root directory. The return value is a `struct dentry` pointer which can be used to create files in this directory. The pointer is also used to clean up the directory in the end, when the device driver is unloaded [8].

²For more information about the `sysfs` file system [6] gives a nice outline

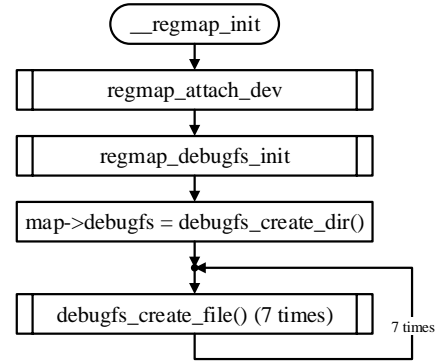


Fig. 4. Initialisation of the debugging structure

To create files in the `debugfs` directory, the most common way is to use the function

```
struct dentry *debugfs_create_file()
```

The first transferred parameter is `const char *name` which is the name of the created file. The next parameter is `umode_t mode` to set the access permissions the file should have³. Next, `struct dentry *parent` indicates the directory in which the file should be created. `void *data` will be stored in the `i_private` field of the resulting inode structure. Finally, `const struct file_operations *fops` is used to set the functions for the file operations like the read and write functions [8].

B. Use of debugfs in the regmap structure

Figure 4 shows the initialisation process for the debugging mode in the regmap structure. To manage the initialisation process, the function

```
regmap_debugfs_init()
```

is called. Firstly, the directory is created with the above mentioned function. Next, the initialisation function creates the following eight files with the help of `debugfs_create_file()`.

- name
- registers
- rbtree
- range
- access
- cache_dirty
- cache_only
- cache_bypass

On every file the user has read access permissions. When the `name` file is read, the kernel module returns the driver name `map->dev->driver->name` to the user space. The `registers` file returns all register addresses with the depending register values. This file can be set to writeable if the constant `REGMAP_ALLOW_WRITE_DEBUGFS` is defined. Then, it is possible to modify the register values from user space. The `rbtree` file is read-only and contains information about the used cache type. In this example a red black tree is used and the virtual file returns information about the number of

³This is done by using the Linux file system permissions

used registers and the number of used bytes. The next file, *range*, is read-only and returns the used register ranges from the devices to the user space. By reading the *access* file, the *debugfs* returns the behaviours for every register. The structure is shown in table I.

TABLE I. STRUCTURE OF THE REGISTER PERMISSION ACCESSES

address	readable	writable	volatile	precious
hex-value	Y/N	Y/N	Y/N	Y/N

When reading *cache_dirty*, the file returns Y if the data in the cache is newer than the data on the hardware. Otherwise N will be returned. The last two files are both readable and writable, so it is possible to change the settings from the user space. The first setting is the *cache_only* bit from the regmap struct. If this bit is set, only the cache is modified and not the hardware when a write access is executed. The second setting controls the *cache_bypass*, which comes also from the regmap struct and was mentioned in section II-C.

The read and write accesses use different functions depending on the functionality. But nevertheless, they all use the functions *copy_to_user* and *copy_from_user* to establish the communication between the user space and the kernel.

C. Using the debugging mode

By default all the files, which are created by *regmap_debugfs_init*, are located in the directory */sys/kernel/debug/regmap/spi0.0/*. To read one of the virtual files, you can easily use the essential program *cat*. Write access to the kernel module can be obtained by using *echo* to write into the debugfs files. *Echo* is installed on every linux system.

Nevertheless, it is often useful to write a program to get more opportunities for the communication with the kernel module. With the help of the file operation functions in table II it is possible to achieve this goal and to write a smart user space program. One example is the setting of the stream position pointer. This pointer can be used to indicate to a desired register where data should be written in.

TABLE II. FILE OPERATION TO GET ACCESS TO A KERNEL MODULE OVER A VIRTUAL FILE SYSTEM [10]

Function	Description
fopen	Open a file
fseek	Reposition stream position indicator
fread	Read block of data from stream
fscanf	Read formatted data from stream
fgetc	Get character from stream
fwrite	Write of data stream
fprintf	Write formatted data to stream
fputc	Write character to stream

IV. EXAMPLE IN THE AT86RF230 MODULE

In this section an example is shown, where the regmap structure is used. The example is the *at86rf230.c* module⁴,

⁴The AT86RF230 is a Low Power 2.4 GHz Transceiver for ZigBee, IEEE 802.15.4 and 6LoWPAN

which uses regmap for read and write accesses, where it is possible.

First, the regmap configuration is done by the *regmap_config* struct (Listing 2).

Listing 2. at86rf230: regmap configuration

```
static const struct regmap_config at86rf230_regmap_spi_config = {
    .reg_bits = 8,
    .val_bits = 8,
    .write_flag_mask = CMD_REG | CMD_WRITE,
    .read_flag_mask = CMD_REG,
    .cache_type = REGCACHE_RBTREE,
    .max_register = AT86RF2XX_NUMREGS,
    .writable_reg = at86rf230_reg_writable,
    .readable_reg = at86rf230_reg_readable,
    .volatile_reg = at86rf230_reg_volatile,
    .precious_reg = at86rf230_reg_precious,
};
```

In the function *at86rf230_probe()* the regmap structure is initialised with the function mentioned in section II. This is shown in Listing 3.

Listing 3. at86rf230: regmap initialisation

```
static int at86rf230_probe(struct spi_device *spi)
{
    // ...

    lp->regmap = devm_regmap_init_spi(spi, &at86rf230_regmap_spi_config);
    if (IS_ERR(lp->regmap)) {
        rc = PTR_ERR(lp->regmap);
        dev_err(&spi->dev, "Failed to allocate register map: %d\n",
            rc);
        goto free_dev;
    }
    // ...
}
```

To change a single register value in the at86rf230 module, the function *__at86rf230_write* is used. In this function the regmap function *regmap_write()* is called as discussed in section II. The implementation is shown in Listing 4.

Listing 4. at86rf230: write access

```
static inline int
__at86rf230_write(struct at86rf230_local *lp,
    unsigned int addr, unsigned int data)
{
    bool sleep = lp->sleep;
    int ret;

    // ...

    ret = regmap_write(lp->regmap, addr, data);

    // ...

    return ret;
}
```

In a similar way, the read access is implemented with the function *__at86rf230_read()*.

Finally, Listing 5 shows an example for the register behaviours. In this example the implementation of the readable function is shown. It returns true if the register is a readable one.

Listing 5. at86rf230: implementation of the readable function

```
static bool
at86rf230_reg_readable(struct device *dev, unsigned int reg)
{
    bool rc;

    /* all writable are also readable */
    rc = at86rf230_reg_writable(dev, reg);
    if (rc)
        return rc;

    /* readonly regs */
    switch (reg) {
        case RG_TRX_STATUS:
        case RG_PHY_RSSI:
        case RG_IRQ_STATUS:
        case RG_PART_NUM:
```

```

    case RG_VERSION_NUM:
    case RG_MAN_ID_1:
    case RG_MAN_ID_0:
        return true;
    default:
        return false;
}

```

While regmap gives a good tool to factor out duplicated code, it is sometimes necessary to use basic SPI functions without the intermediate regmap layer. One notable example is when completion functions are necessary. There, the regmap structure cannot be used and the common way over the *spi_async* function has to be established. Furthermore, the regmap structure cannot be used if it is necessary to read or write whole data blocks and not only one register. An example of the implementation with the *spi_async* function is shown in listing 6. In this example *at86rf230_write_frame_complete* is set as the callback function and *spi_async* is used to write into the frame buffer.

Listing 6. at86rf230: write access with spi_async

```

static void
at86rf230_write_frame(void *context)
{
    struct at86rf230_state_change *ctx = context;
    struct at86rf230_local *lp = ctx->lp;
    struct sk_buff *skb = lp->tx_skb;
    u8 *buf = ctx->buf;
    int rc;

    lp->is_tx = 1;

    buf[0] = CMD_FB | CMD_WRITE;
    buf[1] = skb->len + 2;
    memcpy(buf + 2, skb->data, skb->len);
    ctx->tx.len = skb->len + 2;
    ctx->msg.complete = at86rf230_write_frame_complete;
    rc = spi_async(lp->spi, &ctx->msg);
    if (rc) {
        ctx->tx.len = 2;
        at86rf230_async_error(lp, ctx, rc);
    }
}

```

All code examples about the *at86rf230* module are taken from [11] and are from the kernel version 4.6.

V. CONCLUSION

This paper provides a short overview about the basic implementations and functionalities of the regmap subsystem. It was explained that the use of the regmap subsystem can reduce code doubling and SPI traffic. Whenever only one register value has to be read or written the regmap structure can simplify these accesses. Nevertheless, when a callback function is needed or a whole frame has to be read/write regmap cannot be used.

Another helpful feature of the regmap subsystem, the debugging functionality, was been presented. The possibility to have access to the register values while programming a kernel module can be very helpful.

REFERENCES

- [1] ASoC, <http://www.alsa-project.org/main/index.php/ASoC>, 2016.
- [2] *Linux Source Tree*, lxr.free-electrons.com/source/include/linux/regmap.h, 2016.
- [3] G. Ascheid, *Grundgebiete der Informatik 3* Institute for Communication nologies and Embedded Systems, RWTH Aachen, 2013.
- [4] L. Erdődi, *File compression with LZO algorithm using NVIDIA CUDA architecture* buda University, Faculty of John von Neumann, Budapest, Hungary, 2012.

- [5] *Linux Source Tree*, lxr.free-electrons.com/source/drivers/base/regmap/internal.h, 2016.
- [6] P. Mochel, *The sysfs Filesystem (an extract from: Proceedings of the Linux Symposium)*, Volume One. Ottawa, Ontario, Canada, July 20nd to 23th, 2005
- [7] corbet, *Debugfs*, <http://lwn.net/Articles/115405/>, December 13, 2004
- [8] *Documentation debugfs*, <http://www.mjmwired.net/kernel/Documentation/filesystems/debugfs.txt>, February 11, 2015
- [9] *ubuntuusers*, <https://wiki.ubuntuusers.de/>, 2016
- [10] *cplusplus*, <http://www.cplusplus.com/reference/cstdio/>, 2016
- [11] *Linux Source Tree*, <http://lxr.free-electrons.com/source/drivers/net/ieee802154/at86rf230.c>, 2016